

MAS_4

November 29, 2020

Leonard Vorbeck

No Group

Student Number : 2709813

```
[34]: import pandas as pd
import numpy as np
from scipy.stats import pearsonr as r
import matplotlib.pyplot as plt
from scipy.stats import norm, uniform
%config InlineBackend.figure_format = 'retina'
plt.rcParams["figure.dpi"] = 93
plt.rcParams["figure.figsize"] = (12,7)
plt.style.use("default")
```

0.0.1 4.1.1

For i.i.d sample $X_1, X_2, \dots, X_n \sim N(0, 1)$ with pdf f we have

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$$

with

$$E[\bar{X}_n] = \mu = 0$$

and

$$Var[\bar{X}_n] = \frac{\sigma^2}{n} = \frac{1}{n}$$

For $g(x) = \cos^2(x)$ the MC-estimate with samples drawn from f is

$$E[g(X)] \approx \tilde{g}_n(x) = \frac{1}{n} \sum_{i=1}^n g(x_i)$$

However, the MC-estimator is a random variable itself :

$$\tilde{g}_n(X) \sim N(\mu, \frac{\sigma^2}{n})$$

The uncertainty of the MC-estimator is given by its variance:

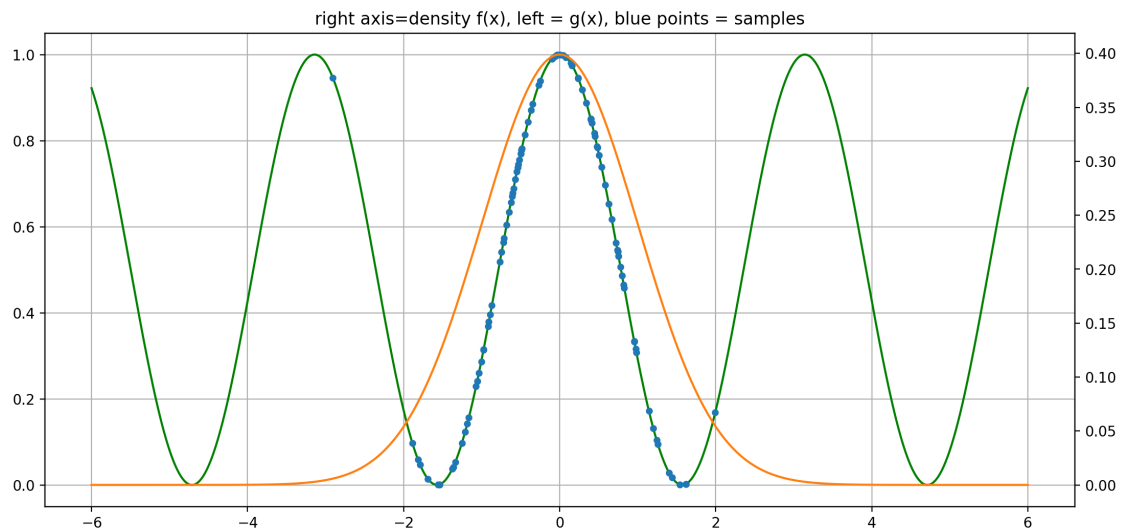
$$\text{Var}(\tilde{g}_n(X)) = \text{Var}[E[g(X)]] \approx \frac{\sigma^2}{n}$$

```
[2]: # Define function
def g(x): return np.cos(x)**2
# Instantiate density
N = norm(0, np.sqrt(1))
```

```
[36]: n = 100
draws = pd.Series(N.rvs(n))
sim = pd.Series([g(x) for x in draws.values], index=draws.values)
func = pd.Series([g(x) for x in np.linspace(-6, 6, 10_000)],
                  index=np.linspace(-6, 6, 10_000))
pdf_x = pd.Series([N.pdf(x) for x in np.linspace(-6, 6, 10_000)],
                  index=np.linspace(-6, 6, 10_000))
```

```
[37]: ax = func.plot(color="green",figsize=(13,6))
sim.plot(ax=ax, style=".",ms=7.73)
pdf_x.plot(ax=ax, secondary_y=True,grid=True,title="right axis=density f(x), left = g(x), blue points = samples")
```

```
[37]: <AxesSubplot:label='be37b585-0741-4cc6-9e12-66495072af7c'>
```



```
[21]: sim.mean(), sim.var()
```

```
[21]: (0.5642320564431826, 0.11467367026497269)
```

The blue dots represent samples $g(X_i)$. The orange line is the original density and the green line is $g(x)$. The chart shows nicely why the MC-estimate is sharp for large values of n .

The simulation reveals $E[g(x)] = 0.56$. The estimate has a variance of $Var[g(x)] = 0.12$.

A better way of estimating uncertainty is to look at the variance of the *estimator* and not the variance of the *estimate*.

```
[27]: n_draws = 100
n_sims = 50
sims = pd.Series([pd.Series([g(x) for x in N.rvs(n_draws)]).mean() for n_sim in
    range(n_sims)])
```

Repeating the experiment 50 times ($n = 100$ each) yields the distribution of the MC-estimator, with

$$E[E[g(x)]] = 0.5675 = E[g(x)] = 0.56 \approx \mu Var[E[g(x)]] = 0.0012 = \frac{Var[g(x)]}{n} = \frac{0.12}{100} \approx \frac{\sigma^2}{n}$$

```
[28]: sims.mean(), sims.var() # For 50 simulation with n=100 each [MC-ESTIMATOR DIST]
```

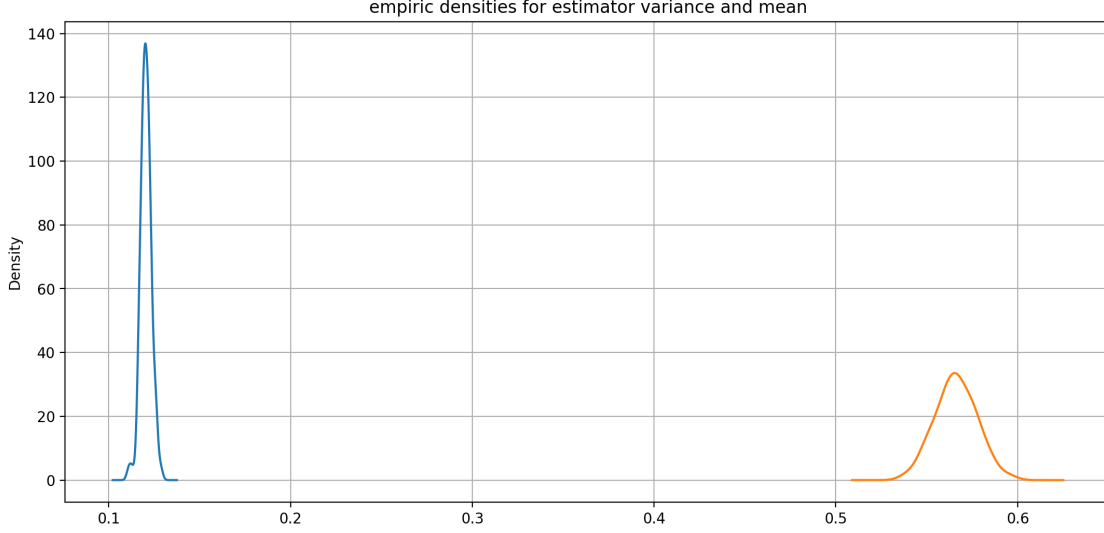
```
[28]: (0.5595546837689471, 0.000924285604112229)
```

```
[29]: sim.mean(), sim.var() # Single simulation with n=100 [ESTIMATE DIST]
```

```
[29]: (0.5642320564431826, 0.11467367026497269)
```

```
[30]: # CLT Induce, blue : VAR[g(x)], orange : AVG[g(x)]
ax = pd.Series([pd.Series([g(np.random.normal(0,1)) for _ in range(1000)]).
    var() for _ in range(100)]).plot.density()
pd.Series([pd.Series([g(np.random.normal(0,1)) for _ in range(1000)]).mean()
    for _ in range(100)]).plot.density(ax=ax,
    grid=True,figsize=(13,6),
    title="empiric densities for estimator
    variance and mean")
```

```
[30]: <AxesSubplot:title={'center': 'empiric densities for estimator variance and
mean'}, ylabel='Density'>
```



0.0.2 4.1.2

For the plain vanilla p-test we have $df = 10 - 2 = 8$ and critical values of -0.632 and 0.632 , which translates to insignificant values for $r \in [-0.632, 0.632]$ and therefore also for $r = 0.3$. Note that p-values are never useful other than for synthetic data where the assumptions of the significance test are met.

Using MC, we can setup a simulation to evaluate how random the found result might be, without assuming the shapes of f_S and f_A .

Let $A \leftarrow [a_1, a_2, \dots, a_{10}]$ and $S \leftarrow [s_1, s_2, \dots, s_{10}]$ with linear correlation $r(A, S)$

Now to check if $r(A, S)$ is non-random (i.e significant), simulate the initial experiment n times to obtain N different values for $r(A, S)$, where in each of the N simulations, S (or both A and S) is permuted ($s_i^t \neq s_i^{t+1}$).

Now after obtaining N values (r_1, r_2, \dots, r_N) , for a one-tailed test sort the correlations, for a two tailed test sort the absolute correlations. The result is a sorted set R indexed i . Now the c CI level corresponds to the $N \times c$ -th element of R , which is $r_{N \times c}$, e.g, r_{990} for $N = 1000$ and $c = .99$.

0.0.3 4.2.1

Let $\varphi(X) = X^2$ and $X \sim N(0, 1)$ with density f and $g \sim U(-5, 5)$ with density g . Then

$$E_f[\varphi(X)] = E_g \left[\varphi(X) \frac{f(X)}{g(X)} \right] \approx \frac{1}{n} \sum_{i=1}^n \varphi(X_i) \frac{f(X_i)}{g(X_i)}$$

where X_i is drawn from g .

First of all, $g(X_i)$ will never be 0, while $f(X_i)$ gets very close to zero for unlikely values, especially if $|X_i| > 1.96$.

In particular, any X_i drawn from U yields $g(X_i) = \frac{1}{10}$. Therefore likelihood ratio can be simplified

$$\frac{f(X_i)}{g(X_i)} \leftarrow f(X_i) \times 10$$

Therefore I predict the following *estimate*.

$$E_f[\varphi(X)] \approx \frac{1}{n} \sum_{i=1}^n 10\varphi(X_i) f(X_i)$$

and

$$VAR_f[\varphi(X)] \approx \frac{1}{n} \sum_{i=1}^n [10\varphi(X_i) f(X_i) - E_f[\varphi(X)]]^2$$

And subsequently the *estimator*

$$E[E_f[\varphi(X)]] = \mu Var[E_f[\varphi(X)]] = \frac{\sigma^2}{n}$$

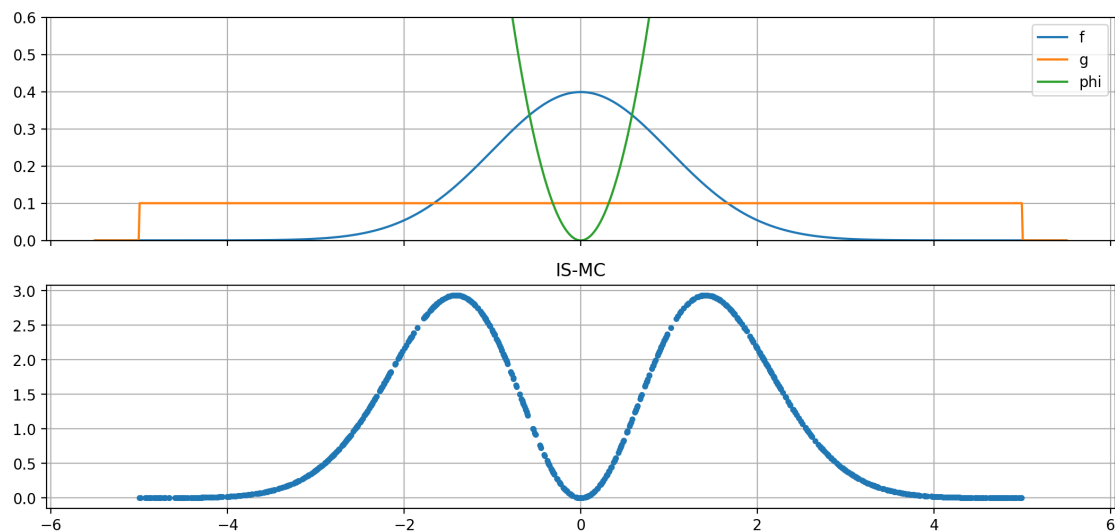
```
[39]: # Instanciate densities
N = norm(0, np.sqrt(1))
U = uniform(loc=-5, scale=10)
```

```
[40]: n_draws = 1000
draws = pd.Series(U.rvs(n_draws))
sim = pd.Series([(xi**2)*10*N.pdf(xi) for xi in draws], index=draws)
sim.mean(), sim.var()
```

```
[40]: (1.0327807495961905, 1.1611606847404532)
```

```
[41]: fig, axes = plt.subplots(ncols=1, nrows=2, figsize=(15,6), sharex=True)
df = pd.DataFrame({"f" : [N.pdf(x) for x in np.linspace(-5.5,5.5,1000)],
                  "g" : [U.pdf(x) for x in np.linspace(-5.5,5.5,1000)],
                  "phi" : [ x**2 for x in np.linspace(-5.5,5.5,1000)]},
                  index=np.linspace(-5.5,5.5,1000))
ax = df.plot(grid=True, figsize=(13,6), ylim=(0,.6), ax=axes[0])
sim.plot(style=".", ax=axes[1], grid=True)
axes[1].set_title("IS-MC")
```

```
[41]: Text(0.5, 1.0, 'IS-MC')
```



The IS-MC-estimate is $E_f[X^2] \approx 1$, with $VAR_f[X^2] \approx 1.05$. (Please note !!! : $E_f[X^2] = E_g\left[X^2 \frac{f(X)}{g(X)}\right]$, similar to VAR, as described earlier).

```
[42]: n_draws = 1000
      n_sims = 50
      sims = pd.Series([pd.Series([(xi**2)*10*N.pdf(xi) for xi in U.rvs(n_draws)]).
        ↳mean() for n_sim in range(n_sims)])
```

If we look at the distribution of the IS-MC-Estimator, for example 50 simulations, each $n = 100$ draws, we get :

```
[43]: sims.mean(), sims.var() # For 50 simulation with n=1000 each [IS-MC-ESTIMATOR,
        ↳DIST]
```

```
[43]: (1.004474884074828, 0.001067546583881102)
```

```
[44]: sim.mean(), sim.var() # For single sim with n=1000 [IS-MC-ESTIMATE DIST]
```

```
[44]: (1.0327807495961905, 1.1611606847404532)
```

```
[45]: sims.var() * n_draws # Fits theoretical Approximation
```

```
[45]: 1.067546583881102
```

Repeating the experiment 50 times ($n = 1000$ each) yields the distribution of the MC-estimator, with

$$E[E[g(x)]] = 1 = E[g(x)] = 1 \approx \mu Var[E[g(x)]] = 0.00099 = \frac{Var[g(x)]}{n} = \frac{1}{1000} \approx \frac{\sigma^2}{n}$$

0.0.4 4.2.2

From the exercise it is not clear whether f_X is the target or proposal distribution. I assume that you meant f_X as the target dist using again U to draw from. In this case, we should use

$$g \sim U(-1, 1)$$

because f_X is not defined for values outside of $[-1, 1]$.

Let $\varphi(X) = X^2$ and $x \in [-1, 1]$ with target density $f(x) = \frac{1+\cos(\pi x)}{2}$ and proposal density $g \sim U(-1, 1)$.

Again

$$E_f[\varphi(X)] = E_g \left[\varphi(X) \frac{f(X)}{g(X)} \right] \approx \frac{1}{n} \sum_{i=1}^n \varphi(X_i) \frac{f(X_i)}{g(X_i)}$$

For $g \sim (-1, 1)$ we have

$$\frac{f(X_i)}{g(X_i)} \leftarrow f(X_i) \times 2$$

s.t.

$$E_f[\varphi(X)] \approx \frac{1}{n} \sum_{i=1}^n 2\varphi(X_i) f(X_i)$$

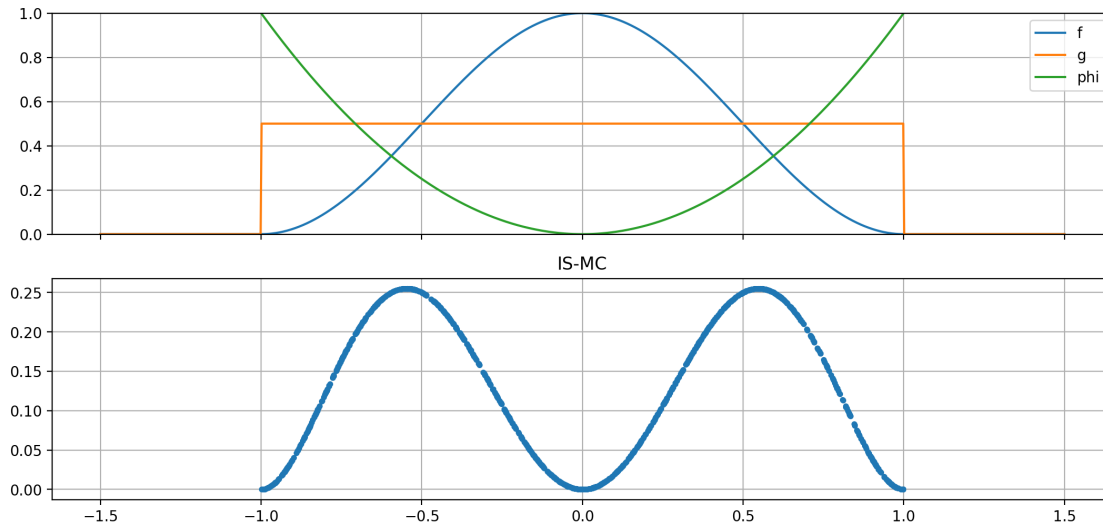
```
[46]: # Define new target density
def fx(x) : return (1+np.cos(np.pi*x))/2 if -1 <= x <= 1 else 0
# Instanciate new proposal density
U = uniform(loc=-1, scale=2)
```

```
[47]: n_draws = 1000
draws = pd.Series(U.rvs(n_draws))
sim = pd.Series([(xi**2)*2*fx(xi) for xi in draws], index=draws)
sim.mean(), sim.var()
```

```
[47]: (0.12879984635877567, 0.0083807124088555)
```

```
[48]: fig, axes = plt.subplots(ncols=1, nrows=2, figsize=(15,6), sharex=True)
df = pd.DataFrame({"f" : [fx(x) for x in np.linspace(-1.5,1.5,1000)],
                  "g" : [U.pdf(x) for x in np.linspace(-1.5,1.5,1000)],
                  "phi" : [ x**2 for x in np.linspace(-1.5,1.5,1000)]},
                  index=np.linspace(-1.5,1.5,1000))
ax = df.plot(grid=True, figsize=(13,6), ylim=(0,1), ax=axes[0])
sim.plot(style=".", ax=axes[1], grid=True)
axes[1].set_title("IS-MC")
```

[48]: Text(0.5, 1.0, 'IS-MC')



The IS-MC-estimate is $E_f[X^2] \approx 0.13$, with $VAR_f[X^2] \approx 0.008$. (Please note !!! : $E_f[X^2] = E_g \left[X^2 \frac{f(X)}{g(X)} \right]$, similar to VAR, as described earlier).

```
[49]: n_draws = 1000
      n_sims = 50
      sims = pd.Series([pd.Series([(xi**2)*2*fx(xi) for xi in U.rvs(n_draws)]).mean(),
      ↪ for n_sim in range(n_sims)])
```

If we look at the distribution of the IS-MC-Estimator, for example 50 simulations, each $n = 1000$ draws, we get :

```
[52]: sims.mean(), sims.var() # For 50 simulation with n=1000 each [IS-MC-ESTIMATOR,
      ↪ DIST]
```

[52]: (0.13101941825232205, 6.673497253745432e-06)

```
[53]: sim.mean(), sim.var() # For single sim with n=1000 [IS-MC-ESTIMATE DIST]
```

[53]: (0.12879984635877567, 0.0083807124088555)

```
[54]: sims.var() * n_draws # Fits theoretical Approximation
```

[54]: 0.006673497253745432

Repeating the experiment 50 times ($n = 1000$ each) yields the distribution of the MC-estimator, with

$$E[E[g(x)]] = 0.13 = E[g(x)] = 0.13 \approx \mu \text{Var}[E[g(x)]] = 9.322359799403895e^{-06} = \frac{\text{Var}[g(x)]}{n} = \frac{0.008}{1000} \approx \frac{\sigma^2}{n}$$

0.0.5 4.3.1

Let

$$KL(f, g) = \int_{-\infty}^{\infty} f(x) \log \frac{f(x)}{g(x)} dx \quad f \sim N(\mu, \sigma^2) \text{ and } g \sim N(\nu, \tau^2)$$

plug the two univariate gaussians into the equation.

$$\dots \dots KL(f, g) = \log\left(\frac{\tau}{\sigma}\right) - \frac{1}{2} + [\sigma^2 + (\mu - \nu)^2] \times 2\tau^{-2}$$

yields something like that (derivation in latex is impossible..)

Obviously KL is not symmetric ($KL(f, g) \neq KL(g, f)$), $\log(\frac{\tau}{\sigma})$ swaps coefficients

```
[220]: np.log((b/a))
```

```
[220]: 0.49247648509779424
```

```
[55]: # In code
def KL(p, q): return np.log(q.std()/p.std()) + (p.var() + (p.mean() - q.
    ↪mean())**2)/(2*q.var()) - (1/2)
```

0.0.6 4.3.2

MC-Estimators

$$\hat{KL}(f, g) = \frac{1}{n} \sum_{i \sim f} \log \frac{f(x_i)}{g(x_i)}$$

$$\hat{KL}(g, f) = \frac{1}{n} \sum_{i \sim g} \log \frac{g(x_i)}{f(x_i)}$$

Let

$$f_X \leftarrow N(10, 2) \quad g_X \leftarrow N(1, 3)$$

```
[56]: def KL(p, q): return np.log(q.std()/p.std()) + (p.var() + (p.mean() - q.
    ↪mean())**2)/(2*q.var()) - (1/2)
```

```
[57]: # Instantiate densities
f = norm(10, np.sqrt(2))
g = norm(1, np.sqrt(3))
```

```
[58]: # MC estimate for KL(f, g)
estimate = []
n=5000
```

```

for xi in f.rvs(n):
    f_xi = f.pdf(xi)
    g_xi = g.pdf(xi)
    res = np.log(f_xi / g_xi)
    estimate.append(res)
pd.Series(estimate).mean()

```

[58]: 13.554201304736548

The MC-estimate of $KL(f, g) = \hat{KL}(f, g) = 13.548$. This is consistent when using the previously derived formula.

[59]: `KL(f, g)`

[59]: 13.536065887387418

```

[60]: # KL(g, f)
estimate = []
n=5000
for xi in g.rvs(n):
    f_xi = f.pdf(xi)
    g_xi = g.pdf(xi)
    res = np.log(g_xi / f_xi)
    estimate.append(res)
pd.Series(estimate).mean()

```

[60]: 20.155027335922334

[61]: `KL(g, f)`

[61]: 20.297267445945913

Again, $KL(g, f) = \hat{KL}(g, f)$

0.0.7 5.1 k -armed-bandit

Let

$$L(t) = E \left(\sum_{i=1}^t (q^* - q(a_i)) \right)$$

be the total (cumulative) regret at time t .

Let

$$Q_t(a) \simeq E[R_t \mid A_t = a] = \frac{\sum_{i=1}^{N_a} r_i}{N_a}$$

be the estimate for action a at time t . (i.e. the sample mean).

with update rule that is equivalent to sample-mean-sharpening (i.e linear)

$$Q_{t+1}(a) = Q_t(a) + \frac{1}{N_a + 1} [r_{t+1} - Q_t(a)]$$

0.0.8 5.1.0 Theoretical optimal lower bound

$$L(t) \geq A \log(t) \quad \text{where} \quad A = \sum_{a: \Delta_a \neq 0} \frac{\Delta_a}{KL(f_a || f_a^*)} \quad \text{and} \quad \Delta_a = q^* - q(a)$$

```
[83]: def L(t, k_dists):
    q_star = [_ for _ in k_dists if _.mean() == max([_.mean() for _ in
↪k_dists])][0]
    A = sum([(q_star.mean() - q_a.mean())/KL(q_a, q_star) for q_a in k_dists if
↪q_a != q_star])
    return np.log(t) * A
```

0.0.9 5.1.1 ϵ -greedy

Action probabilities are handy in this approach :

$$P(a^*) = (1 - \epsilon) \text{ for } a^* = \arg \max_a Q_t(a)$$

$$P(A_{-a^*}) = \epsilon$$

```
[63]: def greedy_bandit(e,k, init_dists, T,
    opt_init=1.05):
    k_dists = init_dists
    q_star = [_ for _ in k_dists if _.mean() == max([_.mean() for _ in
↪k_dists]) ][0]
    q_init = q_star.mean()*opt_init

    # Placeholder for Q values
    Q = { t : {a : q_init if t == 0 else None for a in range(k)} for t in
↪range(T) }

    # Placeholder for action choices
    N = { a : 1 for a in range(k)}
    # Placeholder for regrets
    L = { t : None for t in range(T)}
    # Placeholder for rewards
    R = { t : None for t in range(T)}

    for t in range(T-1):
        ## Choice policy
```

```

    argmax_a = [a for a in Q[t] if Q[t][a] == max([Q[t][a] for a in
↪Q[t]])][0]
    A_ = [a for a in Q[t] if a != argmax_a]
    choice = argmax_a if np.random.binomial(1, e) == 0 else np.random.
↪choice(A_)
    # Get reward from chosen density
    reward = k_dists[choice].rvs(1)[0]
    # Upadte chosen action estimate
    Q[t+1][choice] = (Q[t][choice] + (1/(N[choice]+1))*(reward -
↪Q[t][choice] ))
    # Add 1 obs to N_choice
    N[choice] += 1
    # Update all other
    for a in Q[t]:
        if a == choice:
            continue
        Q[t+1][a] = Q[t][a]
    ## Regret
    regret = q_star.mean() - k_dists[choice].mean()
    L[t] = regret
    ## Reward
    R[t] = reward

    return dict(Q=Q, R=R, L=L, N=N, dists=k_dists, q_star=q_star)

```

0.0.10 5.1.2 UCB

UCB is even simpler, with decision rule according to the most promising action according to its upper boundary.

Choose a s.t.

$$\arg \max_a \left(Q_t(a) + c \sqrt{\frac{\log t}{N_t(a)}} \right)$$

```

[153]: def ucb_bandit(c,k, init_dists, T):
    k_dists = init_dists
    q_star = [_ for _ in k_dists if _.mean() == max([_.mean() for _ in
↪k_dists)) ][0]
    q_init = 0
    u_init = 999
    # Placeholder for U values
    #U = { t : {a : u_init if t == 0 else None for a in range(k)} for t in
↪range(T) }
    # Placeholder for Q values

```

```

    Q = { t : {a : q_init if t == 0 else None for a in range(k)} for t in
    ↪range(T) }
    # Placeholder for action choices
    N = { a : 0 for a in range(k)}
    # Placeholder for regrets
    L = { t : None for t in range(T)}
    # Placeholder for rewards
    R = { t : None for t in range(T)}

    for t in range(T-1):
        ## Choice policy
        U = { a : c * np.sqrt((np.log(t)/N[a])) for a in range(k)}
        choice_values = [ U[a] + Q[t][a] for a in range(k)]
        ucb_argmax_a = choice_values.index(max(choice_values))
        choice = ucb_argmax_a
        A_ = [a for a in Q[t] if a != choice]
        # Get reward from chosen density
        reward = k_dists[choice].rvs(1)[0]
        # Upadte chosen action estimate
        Q[t+1][choice] = (Q[t][choice] + (1/(N[choice]+1))*(reward -
    ↪Q[t][choice] ))
        # Add 1 obs to N_choice
        N[choice] += 1
        # Update all other
        for a in Q[t]:
            if a == choice:
                continue
            Q[t+1][a] = Q[t][a]
        ## Regret
        regret = q_star.mean() - k_dists[choice].mean()
        L[t] = regret
        ## Reward
        R[t] = reward

    return dict(Q=Q, R=R, L=L, N=N, dists=k_dists, q_star=q_star)

```

The arm densities are instantiated as follows :

```

[160]: k = 5
        unit_std = 30
        mean_space=list(range(-100,500,30))
        k_dists = [norm(np.random.choice(mean_space),unit_std) for _ in range(k)]
        [(a.mean(), a.std()) for a in k_dists]

```

```

[160]: [(230.0, 30.0), (20.0, 30.0), (380.0, 30.0), (410.0, 30.0), (440.0, 30.0)]

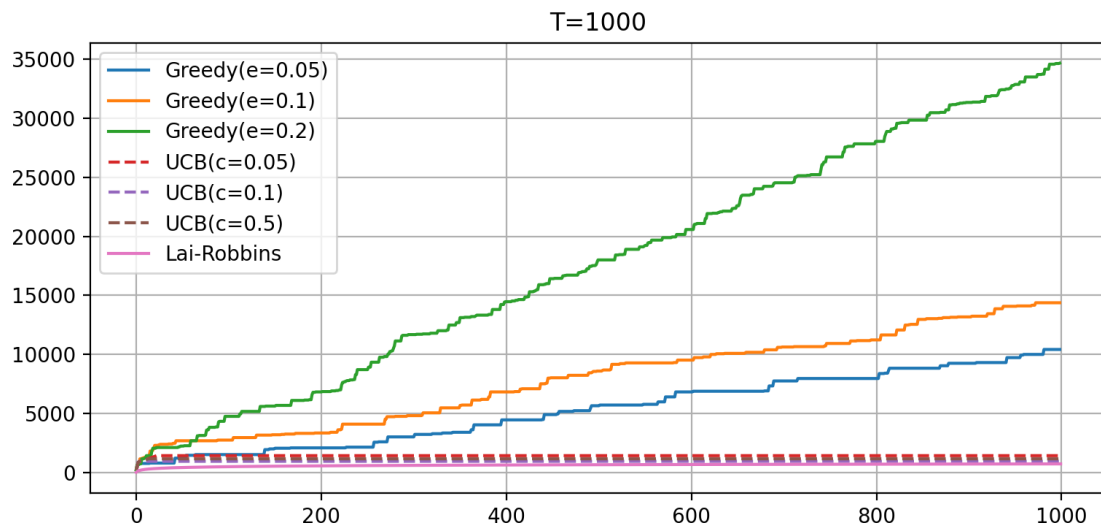
```

```
[163]: e_g = pd.DataFrame({"Greedy(e=%s)" % e :
    ↳ greedy_bandit(e=e,k=k,init_dists=k_dists,T=1000)["L"] for e in [0.05, 0.1, 0.
    ↳ 2]})
ucb = pd.DataFrame({"UCB(c=%s)" % c :
    ↳ ucb_bandit(c=c,k=k,init_dists=k_dists,T=1000)["L"] for c in [0.05, 0.1, 0.
    ↳ 5]})
df = e_g.merge(ucb, right_index=True, left_index=True)
df_cumul = df.cumsum()
df_cumul["Lai-Robbins"] = [L(t+1, k_dists) for t in df.index.values]
correct = df[df == 0].replace(0,1).replace(np.nan, 0)
```

```
/opt/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:19:
RuntimeWarning: divide by zero encountered in log
/opt/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:19:
RuntimeWarning: invalid value encountered in sqrt
/opt/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:19:
RuntimeWarning: invalid value encountered in double_scalars
/opt/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:19:
RuntimeWarning: divide by zero encountered in double_scalars
```

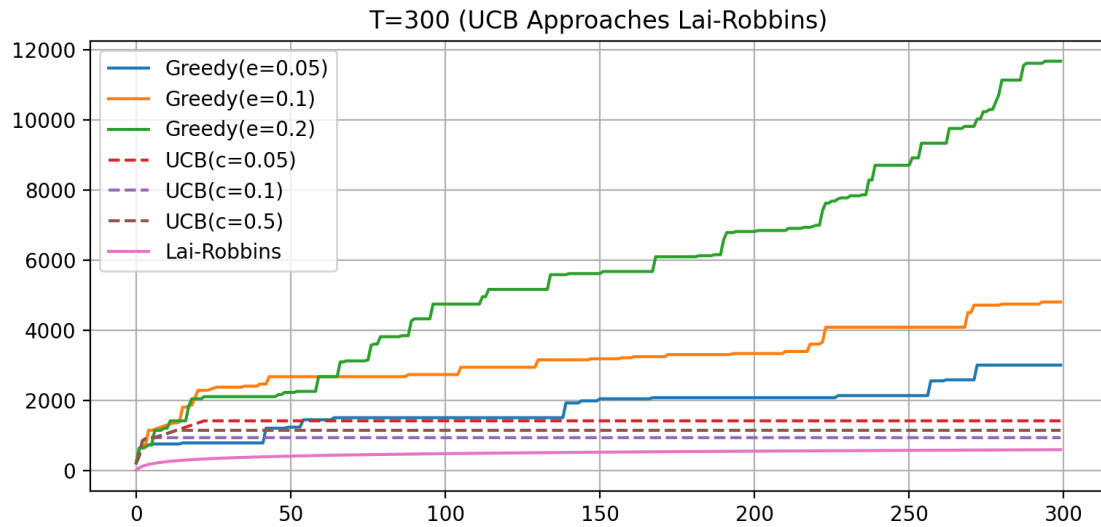
```
[164]: df_cumul.plot(style=["-", "-", "-", "--", "--", "--"],grid=True,
    ↳ figsize=(9,4),title="T=1000")
```

```
[164]: <AxesSubplot:title={'center':'T=1000'}>
```



```
[216]: df_cumul[0:300].plot(style=["-", "-", "-", "--", "--", "--"],grid=True,
    ↳ figsize=(9,4),title="T=300 (UCB Approaches Lai-Robbins)")
```

```
[216]: <AxesSubplot:title={'center':'T=300 (UCB Approaches Lai-Robbins)'}>
```



0.0.11 Correct choices

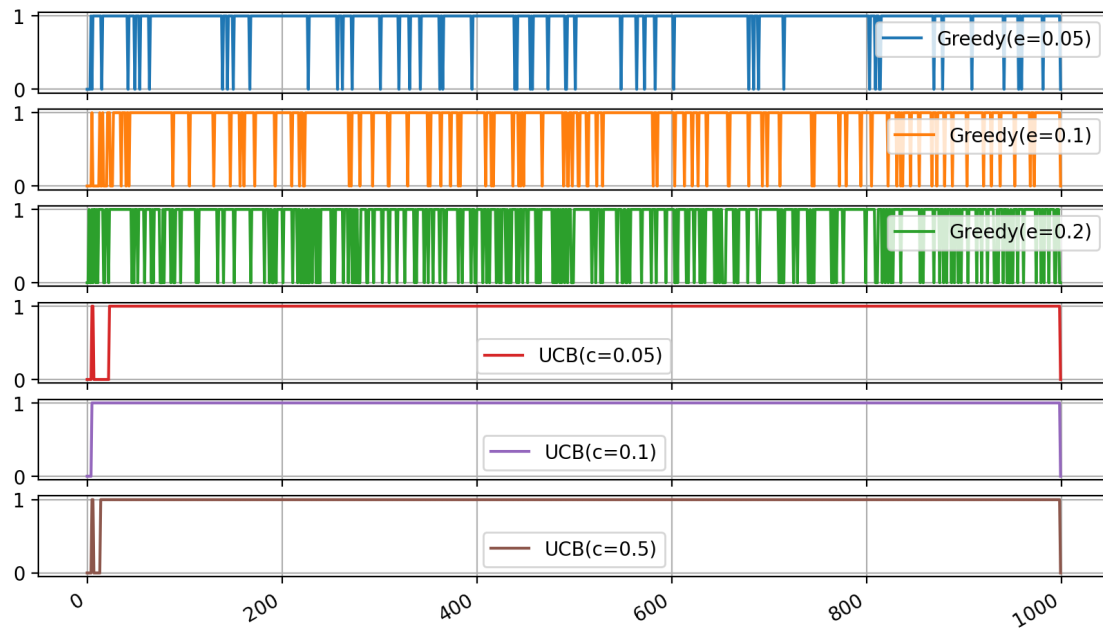
Distribution :

```
[167]: correct.sum()/len(correct)
```

```
[167]: Greedy(e=0.05)    0.946
Greedy(e=0.1)         0.896
Greedy(e=0.2)         0.798
UCB(c=0.05)           0.978
UCB(c=0.1)            0.994
UCB(c=0.5)            0.987
dtype: float64
```

```
[218]: correct.plot(grid=True, style=".-",ms=0.2,figsize=(10,6),subplots=True)
```

```
[218]: array([<AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>,
<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



[]: